

Project Documentation: Neural Network Models with NumPy and TensorFlow

Overview

This project comprises three Jupyter notebooks, each showcasing different neural network architectures for various tasks:

1. **Single Perceptron for Regression**
2. **Two-Layer Neural Network for Binary Classification**
3. **Multi-Layer Neural Network for Multi-Class Classification**

linear_regression_single_perceptron.ipynb

This notebook illustrates a neural network model utilizing a single perceptron for linear regression tasks. Two variants are presented: one with a single input feature and another with two input features.

Contents

1. Single Input Perceptron

- **Data Generation:** Synthetic data is created using `make_regression` from sklearn.
- **Model Implementation:**
 - `initialize_parameters` : Initializes weights and biases.
 - `forward_propagation` : Computes the predicted output.
 - `compute_cost` : Calculates the mean squared error cost.
 - `gradient_descent` : Updates parameters using gradient descent.
 - `nn_model` : Trains the model using the above functions.
- **Visualization:** Plots the regression line and data points.

2. Two Input Perceptron

- **Data Preparation:** Reads and preprocesses the house prices dataset.
- **Model Implementation:** Reuses functions from the single input model.
- **Visualization and Evaluation:** Plots the regression results and calculates RMSE and R² score.

NeuralNet_with_Two_Layers.ipynb

This notebook implements a neural network with one hidden layer for binary classification tasks. The hidden layer can have an arbitrary number of neurons.

Contents

1. **Data Generation:** Synthetic data is created using `make_blobs` from sklearn.
2. **Model Implementation:**
 - `initialize_parameters` : Initializes weights and biases for both layers.
 - `forward_propagation` : Computes the predicted output.
 - `compute_cost` : Calculates the binary cross-entropy loss.
 - `gradient_descent` : Updates parameters using gradient descent.
 - `nn_model` : Trains the model using the above functions.
 - `predict` : Makes predictions using the trained model.
 - `plot_decision_boundary` : Visualizes the decision boundary of the trained model.
3. **Visualization:** Plots decision boundaries for different datasets.
4. **Evaluation:** Calculates and displays the RMSE and R² score.

multi_layer_nn.ipynb

This notebook implements a multi-layer neural network for multi-class classification tasks using the MNIST digits dataset.

Contents

1. **Data Preparation:**
 - Loads and preprocesses the MNIST digits dataset.
 - Scales features using `MinMaxScaler`.
 - Splits the data into training and testing sets.
2. **Model Implementation:**
 - `initialize_parameters` : Initializes weights and biases for each layer.
 - `forward_propagation` : Computes the predicted output using softmax activation for the final layer.
 - `compute_cost` : Calculates the categorical cross-entropy loss.
 - `gradient_descent` : Updates parameters using gradient descent.
 - `learning_rate_decay` : Implements learning rate decay over epochs.

- `nn_model` : Trains the model using the above functions.
- `predict` : Makes predictions using the trained model.

3. Evaluation:

- Evaluates the model using classification report and accuracy score.
- Displays confusion matrix.
- Visualizes misclassified examples.

Neural Network Model with a Single Perceptron and One Input Node

```
In [ ]: # Import necessary libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import tensorflow as tf
from sklearn.datasets import make_regression
```

2024-05-25 20:09:49.694609: I tensorflow/core/platform/cpu_feature_guard.cc:182] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.
To enable the following instructions: SSE4.1 SSE4.2 AVX AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.

```
In [ ]: # Generate synthetic data for regression using sklearn's make_regression
# This creates a dataset with 30 samples and 1 feature with some noise added
X, y = make_regression(n_samples=30, n_features=1, random_state=1, noise=20)

# Transpose X to match the expected input shape for our model
X = X.T
# Reshape y to be a row vector
y = np.reshape(y, (1, -1))

# Print the shapes of X and y to verify their dimensions
print('The shape of X is: ' + str(X.shape))
print('The shape of y is: ' + str(y.shape))
```

The shape of X is: (1, 30)

The shape of y is: (1, 30)

```
In [ ]: def initialize_parameters(n_x, n_y):
    """
    Initialize parameters for the neural network.

    Arguments:
    n_x -- size of the input layer
    n_y -- size of the output layer

    Returns:
    W -- initialized weight matrix of shape (n_y, n_x)
    b -- initialized bias vector of shape (n_y, 1)
    """
    W = tf.Variable(tf.random.normal((n_y, n_x)) * 0.1) # Small random values for weights
    b = tf.Variable(tf.zeros((n_y, 1))) # Biases initialized to zero
    return W, b
```

```
In [ ]: def forward_propagation(X, W, b):
    """
    Perform forward propagation to predict the output.

    Arguments:
    X -- input data of shape (n_x, number of examples)
    W -- weight matrix of shape (n_y, n_x)
    b -- bias vector of shape (n_y, 1)

    Returns:
    y_hat -- predicted output
    """
    y_hat = W @ X + b # Linear combination of inputs and weights plus bias
    return y_hat
```

```
In [ ]: def compute_cost(y, y_hat):
    """
    Compute the cost using mean squared error.

    Arguments:
    y -- true "label" vector
    y_hat -- predicted output vector

    Returns:
    cost -- mean squared error cost
    """
    cost = tf.reduce_mean((y - y_hat) ** 2) / 2 # Mean squared error cost function
    return cost
```

```
In [ ]: def gradient_descent(W, b, dj_dw, dj_db, learning_rate):
    """
```

```

Update parameters using gradient descent.

Arguments:
W -- weight matrix
b -- bias vector
dj_dw -- gradient of the cost with respect to W
dj_db -- gradient of the cost with respect to b
learning_rate -- learning rate for gradient descent

Returns:
W -- updated weight matrix
b -- updated bias vector
"""
W.assign_sub(learning_rate * dj_dw) # Update weights
b.assign_sub(learning_rate * dj_db) # Update biases
return W, b

```

```

In [ ]: def nn_model(X, y, n_x, n_y, epochs, learning_rate, print_cost=True):
        """
        Train the neural network model.

        Arguments:
        X -- input data
        y -- true "label" vector
        n_x -- size of the input layer
        n_y -- size of the output layer
        epochs -- number of epochs to train the model
        learning_rate -- learning rate for gradient descent
        print_cost -- if True, print the cost every 10 epochs

        Returns:
        W -- trained weight matrix
        b -- trained bias vector
        """
        W, b = initialize_parameters(n_x, n_y) # Initialize parameters

        for epoch in range(epochs):
            with tf.GradientTape() as tape:
                y_hat = forward_propagation(X, W, b) # Forward propagation
                cost = compute_cost(y, y_hat) # Compute cost

                if epoch % 10 == 0 and print_cost:
                    print(f'Epoch:{epoch}, Cost: {cost}')

                dj_dw, dj_db = tape.gradient(cost, [W, b]) # Compute gradients
                W, b = gradient_descent(W, b, dj_dw, dj_db, learning_rate) # Update parameters

        W = W.numpy() # Convert TensorFlow variables to NumPy arrays
        b = b.numpy()
        return W, b

```

```

In [ ]: # Set hyperparameters
LEARNING_RATE = 0.05
EPOCHS = 100
n_x = X.shape[0] # Number of input features
n_y = 1 # Number of output features (single output)

# Train the model and get the final parameters
W, b = nn_model(X, y, n_x, n_y, EPOCHS, LEARNING_RATE, print_cost=True)

```

Epoch:0, Cost: 792.9966430664062

2024-05-25 20:09:52.945854: I tensorflow/core/common_runtime/process_util.cc:146] Creating new thread pool with default inter op setting: 2. Tune using inter_op_parallelism_threads for best performance.


Epoch:10, Cost: 370.9805603027344
Epoch:20, Cost: 222.5522918701172
Epoch:30, Cost: 170.1344451904297
Epoch:40, Cost: 151.54661560058594
Epoch:50, Cost: 144.92808532714844
Epoch:60, Cost: 142.5618438720703
Epoch:70, Cost: 141.7124481201172
Epoch:80, Cost: 141.40638732910156
Epoch:90, Cost: 141.29568481445312

```

In [ ]: # Generate a range of x values for plotting the regression line
x = np.linspace(-3, 3, 50)
# Compute the predicted y values using the trained parameters
y_pred = W @ x.reshape(1, -1) + b

# Plot the regression line and the data points
plt.plot(x, y_pred[0], c='r', label='Regression Line')
plt.scatter(X, y, label='Data Points')
plt.xlabel('X')
plt.ylabel('y')
plt.title('Linear Regression with Single Perceptron')
plt.legend()
plt.show()

```



Neural Network Model with a Single Perceptron and Two Input Nodes

```
In [ ]: import seaborn as sns
        from sklearn.model_selection import train_test_split
```

```
In [ ]: # Load house prices dataset
        df = pd.read_csv('house_prices_train.csv', index_col='Id')
```

```
In [ ]: # Check and print the percentage of missing values in each column
        for column in df.columns:
            if np.sum(df[column].isna()) / 1460 > 0:
                print(f'{np.sum(df[column].isna()) / 1460 * 100 :0.2f}% of "{column}" is null.')
```

```
17.74% of "LotFrontage" is null.
93.77% of "Alley" is null.
59.73% of "MasVnrType" is null.
0.55% of "MasVnrArea" is null.
2.53% of "BsmtQual" is null.
2.53% of "BsmtCond" is null.
2.60% of "BsmtExposure" is null.
2.53% of "BsmtFinType1" is null.
2.60% of "BsmtFinType2" is null.
0.07% of "Electrical" is null.
47.26% of "FireplaceQu" is null.
5.55% of "GarageType" is null.
5.55% of "GarageYrBlt" is null.
5.55% of "GarageFinish" is null.
5.55% of "GarageQual" is null.
5.55% of "GarageCond" is null.
99.52% of "PoolQC" is null.
80.75% of "Fence" is null.
96.30% of "MiscFeature" is null.
```

```
In [ ]: # List to store columns with more than 6% missing values
        cols = []
        for column in df.columns:
            if np.sum(df[column].isna()) / 1460 * 100 > 6:
                cols.append(column)

        # Drop columns with more than 6% missing values and rows with any missing values
        df_new = df.drop(columns=cols)
        df_new = df_new.dropna()

        # Print the shapes of the original and cleaned datasets
        print(f'df.shape: {df.shape}, df_new.shape: {df_new.shape}')
```

```
df.shape: (1460, 80), df_new.shape: (1338, 73)
```

```
In [ ]: # Number of features to select for the model
        n_features = 2

        # Convert categorical variables to dummy variables
        df_new = pd.get_dummies(df_new, drop_first=True)


        # Select top correlated features with the target variable 'SalePrice'
        columns = df_new.corrwith(df_new['SalePrice']).abs().nlargest(n_features + 1).keys()[1:]
        columns
```

```
Out[ ]: Index(['OverallQual', 'GrLivArea'], dtype='object')
```

```
In [ ]: # Assign the selected features to X and target variable to y
X = df_new[columns]
y = df_new['SalePrice']

# Visualize the relationship between selected features and target variable
sns.scatterplot(data=df, x='GrLivArea', y=y, hue='OverallQual')
```

```
Out[ ]: <Axes: xlabel='GrLivArea', ylabel='SalePrice'>
```



```
In [ ]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=101)
```

```
In [ ]: # Normalize the features for training and testing sets
X_train_norm = ((X_train - X.mean()) / X.std()).T.to_numpy()
X_test_norm = ((X_test - X.mean()) / X.std()).T.to_numpy()

# Normalize the target variable for training set
y_train_norm = (y_train - y.mean()) / y.std()
y_train_norm = np.reshape(y_train_norm, (1, -1))
y_test = np.reshape(y_test, (1, -1))
```

```
In [ ]: # Set hyperparameters
LEARNING_RATE = 0.05
EPOCHS = 100
n_x = X_train_norm.shape[0] # Number of input features
n_y = 1 # Number of output features (single output)

# Train the model and get the final parameters
W, b = nn_model(X_train_norm, y_train_norm, n_x, n_y, EPOCHS, LEARNING_RATE, print_cost=True)
```

```
Epoch:0, Cost: 0.4375437796115875
Epoch:10, Cost: 0.20974504947662354
Epoch:20, Cost: 0.16759221255779266
Epoch:30, Cost: 0.1587148904800415
Epoch:40, Cost: 0.15617524087429047
Epoch:50, Cost: 0.15507946908473969
Epoch:60, Cost: 0.15445576608181
Epoch:70, Cost: 0.154059499502182
Epoch:80, Cost: 0.15379902720451355
Epoch:90, Cost: 0.15362612903118134
```


```
In [ ]: # Visualize the true data points
sns.scatterplot(data=df, x='GrLivArea', y='SalePrice', hue='OverallQual', palette='Blues', label='True Data')

# Predict the normalized target values using the trained model
y_values_norm = W @ X_test_norm + b
# Convert the normalized predicted values back to original scale
y_values = y_values_norm * y.std() + y.mean()

# Create a dataframe with the test features and predicted target values
df2 = pd.DataFrame(np.hstack([X_test, y_values.T]), columns=[*columns, 'SalePrice'])

# Visualize the predicted data points
sns.scatterplot(data=df2, x='GrLivArea', y='SalePrice', hue='OverallQual', palette='YlOrBr', label='Predicted')
```

```
Out[ ]: <Axes: xlabel='GrLivArea', ylabel='SalePrice'>
```



```
In [ ]: # Evaluate the model using RMSE and R2 score
from sklearn.metrics import mean_squared_error, r2_score

RMSE = np.sqrt(mean_squared_error(y_test.T, y_values.T))
print(f'RMSE : {RMSE}')
print(f'R2 Score: {r2_score(y_test.T, y_values.T)}')
```

RMSE : 40488.32085402417
R2 Score: 0.7128511001638731

Neural Network Model with One Hidden Layer

```
In [ ]: # Import necessary libraries
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
import pandas as pd
from sklearn.datasets import make_blobs
```


```
In [ ]: # Generate synthetic data for classification using sklearn's make_blobs
m = 2000
X, y = make_blobs(m, centers=([2.5, 3], [6.5, 8], [2, 8], [7.5, 3]), random_state=0)
y[(y == 0) | (y == 1)] = 1
y[(y == 2) | (y == 3)] = 0

# Transpose X to match the expected input shape for our model
X = X.T
# Reshape y to be a row vector
y = np.reshape(y, (1, -1))

# Plot the data points
plt.scatter(X[0, :], X[1, :], c=y[0, :], cmap='RdBu', edgecolors='white', linewidths=0.2);

# Print the shapes of X and y to verify their dimensions
print('The shape of X is: ' + str(X.shape))
print('The shape of y is: ' + str(y.shape))
```

The shape of X is: (2, 2000)
The shape of y is: (1, 2000)



```
In [ ]: def initialize_parameters(n_x, n_h, n_y):
    """
    Initialize parameters for the neural network with two layers.

    Arguments:
    n_x -- size of the input layer
    n_h -- size of the hidden layer
    n_y -- size of the output layer

    Returns:
    params -- dictionary containing initialized parameters
    """
    W1 = tf.Variable(tf.random.normal(shape=(n_h, n_x)) * tf.sqrt(2/n_x)) # He initialization for weights
    b1 = tf.Variable(tf.zeros(shape=(n_h, 1))) # Biases initialized to zero
    W2 = tf.Variable(tf.random.normal(shape=(n_y, n_h)) * tf.sqrt(2/n_y)) # He initialization for weights
    b2 = tf.Variable(tf.zeros(shape=(n_y, 1))) # Biases initialized to zero

    params = {
        'W1': W1,
        'b1': b1,
        'W2': W2,
        'b2': b2
    }

    return params
```

```
In [ ]: def sigmoid(z):
    """
    Compute the sigmoid activation function.

    Arguments:
    z -- input to the sigmoid function

    Returns:
    sigmoid(z) -- output of the sigmoid function
    """
    return 1 / (1 + tf.exp(-z))
```

```
In [ ]: def forward_propagation(X, params):
    """
    Perform forward propagation to predict the output.

    Arguments:
    X -- input data of shape (n_x, number of examples)
    params -- dictionary containing initialized parameters

    Returns:
    y_hat -- predicted output
    """
    W1 = params['W1']
    b1 = params['b1']
    W2 = params['W2']
    b2 = params['b2']

    Z1 = W1 @ X + b1 # Linear transformation
    A1 = tf.nn.relu(Z1) # ReLU activation function
    Z2 = W2 @ A1 + b2 # Linear transformation
    y_hat = sigmoid(Z2) # Sigmoid activation function

    return y_hat
```

```
In [ ]: def compute_cost(y, y_hat):
        """
        Compute the cost using binary cross-entropy.

        Arguments:
        y -- true "label" vector
        y_hat -- predicted output vector

        Returns:
        cost -- binary cross-entropy cost
        """
        logloss = tf.keras.losses.binary_crossentropy(y, y_hat) # Binary cross-entropy loss function
        return tf.reduce_mean(logloss)
```

```
In [ ]: def gradient_descent(params, grads, learning_rate):
        """
        Update parameters using gradient descent.

        Arguments:
        params -- dictionary containing parameters
        grads -- dictionary containing gradients of the cost with respect to parameters
        learning_rate -- learning rate for gradient descent

        Returns:
        params -- updated parameters
        """
        for i in params.keys():
            params[i].assign_sub(learning_rate * grads[i]) # Update parameters using gradients

        return params
```

```
In [ ]: def nn_model(X, y, n_x, n_h, n_y, epochs, learning_rate):
        """
        Train the neural network model.

        Arguments:
        X -- input data
        y -- true "label" vector
        n_x -- size of the input layer
        n_h -- size of the hidden layer
        n_y -- size of the output layer
        epochs -- number of epochs to train the model
        learning_rate -- learning rate for gradient descent

        Returns:
        params -- trained parameters
        """
        params = initialize_parameters(n_x, n_h, n_y) # Initialize parameters
        for epoch in range(epochs):
            with tf.GradientTape() as tape:
                y_hat = forward_propagation(X, params) # Forward propagation
                cost = compute_cost(y, y_hat) # Compute cost

                if epoch % 100 == 0:
                    print(f'Epoch:{epoch}, Cost: {cost}')

            grads = tape.gradient(cost, params) # Compute gradients
            params = gradient_descent(params, grads, learning_rate) # Update parameters

        return params
```

```
In [ ]: # Set hyperparameters
LEARNING_RATE = 0.08
EPOCHS = 1000
n_x = X.shape[0] # Number of input features
n_h = 8 # Number of units in hidden layer
n_y = y.shape[0] # Number of output units

# Train the model and get the final parameters
params = nn_model(X, y, n_x, n_h, n_y, EPOCHS, LEARNING_RATE)
```

```
Epoch:0, Cost: 4.5431084632873535
Epoch:100, Cost: 0.34587621688842773
Epoch:200, Cost: 0.291795551776886
Epoch:300, Cost: 0.25851231813430786
Epoch:400, Cost: 0.2341194450855255
Epoch:500, Cost: 0.21590639650821686
Epoch:600, Cost: 0.20175601541996002
Epoch:700, Cost: 0.1903807371854782
Epoch:800, Cost: 0.18064361810684204
Epoch:900, Cost: 0.17172929644584656
```

```
In [ ]: def predict(X, params):
        """
        Make predictions using the trained model.
```



```

Arguments:
X -- input data
params -- trained parameters

Returns:
predictions -- array of predictions
"""
A2 = forward_propagation(X, params) # Forward propagation
predictions = A2 > 0.5 # Convert probabilities to binary predictions

return predictions.numpy()

```

```

In [ ]: def plot_decision_boundary(X, y, params):
        """
        Plot the decision boundary of the trained model.

        Arguments:
        X -- input data
        y -- true "label" vector
        params -- trained parameters
        """

        min1, max1 = X[0, :].min() - 1, X[0, :].max() + 1
        min2, max2 = X[1, :].min() - 1, X[1, :].max() + 1


        # Generate a grid of points within the feature space
        x1grid = np.arange(min1, max1, 0.1)
        x2grid = np.arange(min2, max2, 0.1)
        xx, yy = np.meshgrid(x1grid, x2grid)
        r1, r2 = xx.flatten(), yy.flatten()
        r1, r2 = r1.reshape((1, len(r1))), r2.reshape((1, len(r2)))
        grid = np.vstack((r1, r2))

        # Make predictions on the grid points
        predictions = predict(grid, params)
        zz = predictions.reshape(xx.shape)

        # Plot decision boundary and data points
        plt.contourf(xx, yy, zz, cmap='BrBG')
        plt.scatter(X[0, :], X[1, :], c=y[0, :], cmap='RdBu', edgecolors='white', linewidths=0.2)
        plt.title("Decision Boundary for hidden layer size " + str(n_h));
        plt.xlabel('Feature 1')
        plt.ylabel('Feature 2')
        plt.show()

# Plot decision boundary for the synthetic data
plot_decision_boundary(X, y, params)

```




Additional dataset

```

In [ ]: data = pd.read_csv('Arcs.csv')
X = data.iloc[:, :-1].T.to_numpy()
y = np.reshape(data.iloc[:, -1], (1, -1))
plt.scatter(X[0, :], X[1, :], c=y[0, :], cmap='RdBu', edgecolors='white', linewidths=0.2);

```




```
In [ ]: LEARNING_RATE = 0.5
        EPOCHS = 1000

        n_x = X.shape[0] # Number of input features
        n_h = 8 # Number of units in hidden layer
        n_y = y.shape[0] # Number of output units

        params = nn_model(X, y, n_x, n_h, n_y, EPOCHS, LEARNING_RATE)
```

```
Epoch:0, Cost: 2.0575194358825684
Epoch:100, Cost: 0.2534295320510864
Epoch:200, Cost: 0.20632652938365936
Epoch:300, Cost: 0.1616419553756714
Epoch:400, Cost: 0.13303911685943604
Epoch:500, Cost: 0.11789369583129883
Epoch:600, Cost: 0.10948474705219269
Epoch:700, Cost: 0.10379713773727417
Epoch:800, Cost: 0.09918693453073502
Epoch:900, Cost: 0.09636596590280533
```

```
In [ ]: plot_decision_boundary(X, y, params);
```



Neural Network Model with Multiple Layers

```
In [ ]: # Import necessary libraries
import numpy as np
import pandas as pd
import tensorflow as tf
import matplotlib.pyplot as plt
from sklearn.datasets import load_digits
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
```

```
In [ ]: # Load the digits dataset
digits = load_digits()


# Extract features (X) and labels (y) from the dataset
X = digits['data']
y = digits['target']

# Print the shape of X and y to verify their dimensions
print(f'X.shape: {X.shape}')
print(f'y.shape: {y.shape}')
```

```
X.shape: (1797, 64)
y.shape: (1797,)
```

```
In [ ]: # Visualize a sample image from the dataset
plt.imshow(X[1].reshape(8, 8))
```

```
Out[ ]: <matplotlib.image.AxesImage at 0x7b4bface28d0>
```



```
In [ ]: # Scale the features to the range [0, 1]
scaler = MinMaxScaler()
X = scaler.fit_transform(X)
```

```
In [ ]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=42)
```

```
In [ ]: # Transpose the data to match the expected input shape for our model
X_train = X_train.T
X_test = X_test.T
y_train = tf.keras.utils.to_categorical(y_train).T

# Print the shapes of training and testing data to verify their dimensions
print(f'X_train.shape: {X_train.shape}')
print(f'X_test.shape: {X_test.shape}')
print(f'y_train.shape: {y_train.shape}')
print(f'y_test.shape: {y_test.shape}')
```

```
X_train.shape: (64, 1203)
X_test.shape: (64, 594)
y_train.shape: (10, 1203)
y_test.shape: (594,)
```

```
In [ ]: def initialize_parameters(layer_dims):
    """
    Initialize parameters for the neural network with multiple layers.

    Arguments:
    layer_dims -- list containing the number of units in each layer

    Returns:
    params -- dictionary containing initialized parameters
    """
    params = {}
    for i in range(1, len(layer_dims)):
        params[f'W{i}'] = tf.Variable(tf.random.normal(shape=(layer_dims[i], layer_dims[i-1])) * tf.sqrt(2/layer_dims[i-1]))
        params[f'b{i}'] = tf.Variable(tf.zeros(shape=(layer_dims[i], 1)))

    return params
```

```
In [ ]: def forward_propagation(X, params):
    """
    Perform forward propagation to predict the output.
```

```

Arguments:
X -- input data of shape (n_x, m)
params -- dictionary containing initialized parameters

Returns:
y_hat -- predicted output
"""
l = len(params) // 2
A = X
for i in range(1, l):
    Z = params[f'W{i}'] @ A + params[f'b{i}']
    A = tf.nn.relu(Z)

Z = params[f'W{l}'] @ A + params[f'b{l}']
y_hat = tf.nn.softmax(Z)

return y_hat

```

```

In [ ]: def compute_cost(y, y_hat):
        """
        Compute the cost using categorical cross-entropy.

        Arguments:
        y -- true labels
        y_hat -- predicted probabilities

        Returns:
        cost -- categorical cross-entropy cost
        """
        loss = tf.keras.losses.categorical_crossentropy(y, y_hat)
        return tf.reduce_mean(loss)

```

```

In [ ]: def gradient_descent(params, grads, learning_rate):
        """
        Update parameters using gradient descent.

        Arguments:
        params -- dictionary containing parameters
        grads -- dictionary containing gradients of the cost with respect to parameters
        learning_rate -- learning rate for gradient descent

        Returns:
        params -- updated parameters
        """
        for i in params.keys():
            params[i].assign_sub(learning_rate * grads[i])

        return params

```

```

In [ ]: def learning_rate_decay(learning_rate, epoch_num, decay_rate=1, time_interval=1000):
        """
        Decay the learning rate over time.

        Arguments:
        learning_rate -- initial learning rate
        epoch_num -- current epoch number
        decay_rate -- rate of decay
        time_interval -- time interval for decay

        Returns:
        updated_learning_rate -- decayed learning rate
        """
        updated_learning_rate = learning_rate / (1 + decay_rate * epoch_num / time_interval)
        return updated_learning_rate

```

```

In [ ]: def nn_model(X, y, layer_dims, epochs, learning_rate, decay_rate=1, print_cost=False):
        """
        Train the neural network model.

        Arguments:
        X -- input data
        y -- true labels
        layer_dims -- list containing the number of units in each layer
        epochs -- number of epochs to train the model
        learning_rate -- initial learning rate
        decay_rate -- rate of decay for learning rate
        print_cost -- whether to print the cost during training

        Returns:
        params -- trained parameters
        """
        params = initialize_parameters(layer_dims)
        for epoch in range(epochs):
            with tf.GradientTape() as tape:
                y_hat = forward_propagation(X, params)

```

```

        cost = compute_cost(y, y_hat)

        if print_cost and epoch % 100 == 0:
            print(f'Epoch:{epoch}, Cost: {cost:0.2f}')

        grads = tape.gradient(cost, params)
        decayed_learning_rate = learning_rate_decay(learning_rate, epoch, decay_rate)
        params = gradient_descent(params, grads, decayed_learning_rate)

    return params

```

```

In [ ]: # Set hyperparameters
LEARNING_RATE = 0.025
EPOCHS = 1000
LAYER_DIMS = [X_train.shape[0], 64, 32, y_train.shape[0]] # Number of units in each layer

# Train the model and get the final parameters
params = nn_model(X_train, y_train, LAYER_DIMS, EPOCHS, LEARNING_RATE, decay_rate=1.5, print_cost=True)

```

```

Epoch:0, Cost: 865.67
Epoch:100, Cost: 597.35
Epoch:200, Cost: 581.50
Epoch:300, Cost: 579.44
Epoch:400, Cost: 578.02
Epoch:500, Cost: 577.70
Epoch:600, Cost: 577.15
Epoch:700, Cost: 576.96
Epoch:800, Cost: 576.86
Epoch:900, Cost: 576.76

```

```

In [ ]: def predict(X, params):
        """
        Make predictions using the trained model.

        Arguments:
        X -- input data
        params -- trained parameters

        Returns:
        predictions -- array of predictions
        """
        y_hat = forward_propagation(X, params)
        predictions = np.argmax(y_hat, axis=0)

        return predictions

```

```

In [ ]: # Make predictions using the trained model
y_pred = predict(X_test, params)

```

```

In [ ]: from sklearn.metrics import accuracy_score, classification_report, confusion_matrix, ConfusionMatrixDisplay

# Evaluate the model using classification report and confusion matrix
print(classification_report(y_test, y_pred))
ConfusionMatrixDisplay(confusion_matrix(y_test, y_pred)).plot()


```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	55
1	0.98	1.00	0.99	55
2	0.98	1.00	0.99	52
3	0.98	0.96	0.97	56
4	1.00	0.98	0.99	64
5	0.97	0.99	0.98	73
6	0.98	0.98	0.98	57
7	0.98	0.98	0.98	62
8	0.98	0.98	0.98	52
9	0.99	0.97	0.98	68
accuracy			0.98	594
macro avg	0.99	0.99	0.99	594
weighted avg	0.98	0.98	0.98	594

```

Out[ ]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x7b4c14375950>

```



```
In [ ]: # Make predictions using the trained model
y_pred_train = predict(X_train, params)

# Calculate and print the accuracy score on the training set
accuracy_train = np.mean(y_pred_train == y_train.argmax(axis=0))
print(f'Training Accuracy: {accuracy_train}')
```

Training Accuracy: 1.0

```
In [ ]: # Calculate and print the accuracy score on the test set
accuracy_test = np.mean(y_pred == y_test)
print(f'Test Accuracy: {accuracy_test}')
```

Test Accuracy: 0.9848484848484849

```
In [ ]: def find_closest_factors(number):
    """
    Find the closest factors of a number.

    Arguments:
    number -- input number

    Returns:
    a -- one factor
    b -- another factor
    """
    a = int(np.sqrt(number))
    for i in range(a, 0, -1):
        if number % i == 0:
            return i, number // i
```

```
In [ ]: # Calculate the number of misclassifications
errors = np.sum(y_test != y_pred)

# Find and print the closest factors of the number of misclassifications
a, b = find_closest_factors(errors)
print(f'Number of Misclassifications: {errors} = {a} * {b}')
```

Number of Misclassifications: 9 = 3 * 3

```
In [ ]: # Visualize some of the misclassified digits
# Plot a grid of images of misclassified digits along with their true and predicted labels
misclassified_indices = np.where(y_test != y_pred)[0]
num_misclassified = len(misclassified_indices)

# Create subplots for each misclassified digit
fig, axs = plt.subplots(a, b, figsize=(b*2.75, a*3))
for i, ax in enumerate(axs.flat):
    ax.imshow(X_test[:, misclassified_indices[i]].reshape(8, 8))
    ax.set_title(f'True: {y_test[misclassified_indices[i]]}, Predicted: {y_pred[misclassified_indices[i]]}')

# Show the plot
plt.tight_layout()
plt.show()
```